

Agents: Landscape and Challenges

Manish Sharma
Gen AI Design & Adoption



Agenda



AN OPEN DISCUSSION



(IDEA TO LIVE)


What are Agents & Use Cases



Agents are autonomous systems that perform tasks using reasoning and planning



Identify use cases by
Common domains: IT ticketing, finance reconciliation, document processing



High volume, repetitive workflows

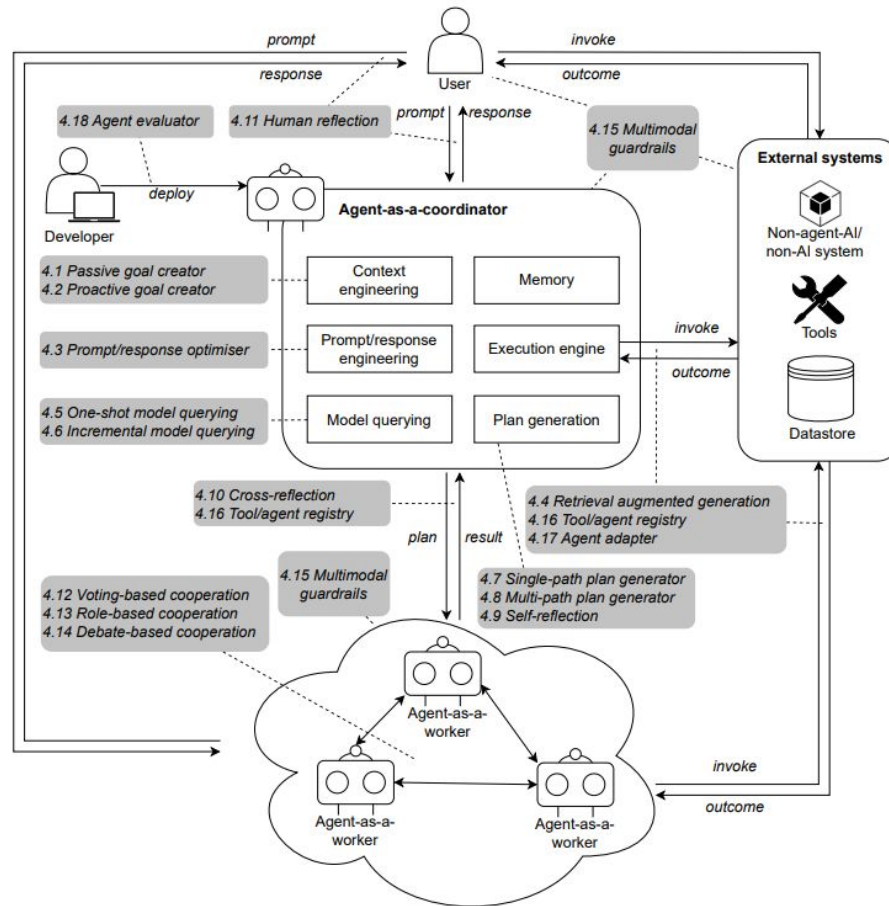


Tasks involving multiple systems




Manual interventions

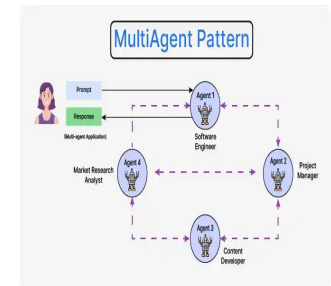
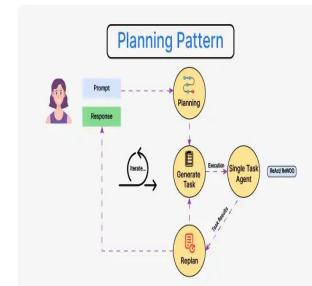
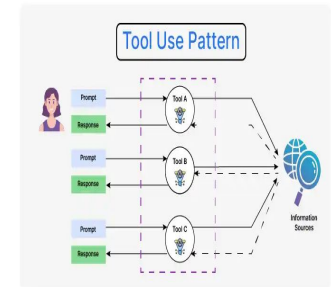
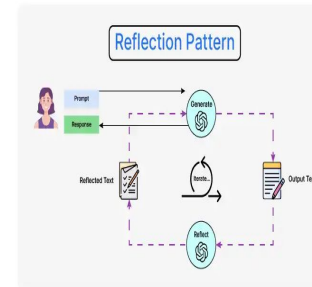
Agents Ecosystem



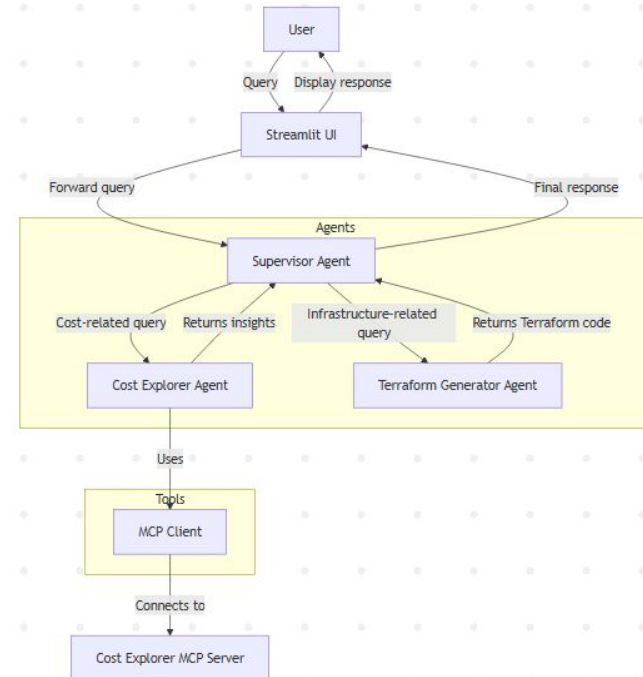
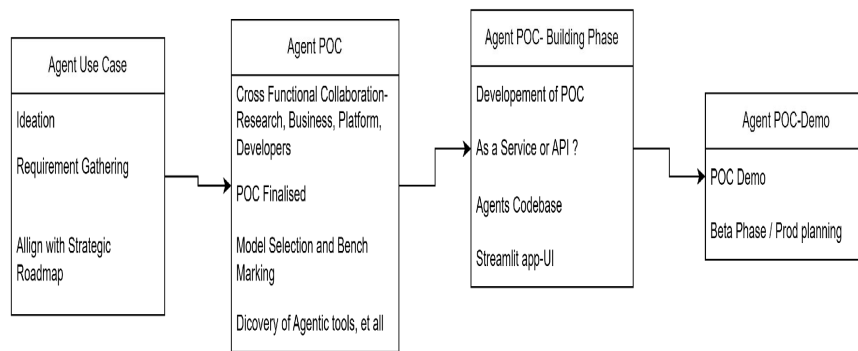
Popular Agent Frameworks



Agentic Design Patterns



Idea to POC



Idea to POC

```
# Supervisor / Orchestrator Agent
# -----
SUPERVISOR_PROMPT = (
    "You are the AI system supervisor. Route user requests to the appropriate specialist.\n"
    "If the user's question is about AWS costs, billing, or spending, use the 'cost_agent' tool. "\n"
    "If it's about generating or understanding AWS infrastructure (e.g. creating resources, Terraform code, architecture), use 'terraform_agent' . "\n"
    "Respond with the chosen tool's output."
)

supervisor_agent = Agent(
    system_prompt=SUPERVISOR_PROMPT,
    tools=[cost_agent, terraform_agent],
    callback_handler=None # suppress internal logs for clarity
)

## Streamlit Chatbot UI

import asyncio
import streamlit as st
from strands import Agent
# (Assume the above agent code is in the same scope or imported)

st.set_page_config(page_title="Multi-Agent AWS Chatbot")
st.header("AWS Multi-Agent Assistant")

# Initialize the supervisor agent (ideally done outside the request loop)
# (Reuse the supervisor_agent defined above)

# Chat history
if "history" not in st.session_state:
    st.session_state.history = []

# User Input
user_input = st.text_input("You:", key="input")

if st.button("Send"):
    prompt = user_input.strip()
    if prompt:
        # Append user message to chat history
        st.session_state.history.append({"role": "user", "text": prompt})

        # Call the supervisor agent (streaming)
        response_text = ""
        placeholder = st.empty()

        # Use Strands streaming interface (asynchronous)
        async def chat_stream():
            try:
                async for event in supervisor_agent.stream_async(prompt):
                    if "data" in event:
                        # Append new chunk to response
                        nonlocal response_text
                        response_text += event["data"]
                        placeholder.markdown(response_text) # update display
            except Exception as e:
                placeholder.text(f"Error: {e}")

        # Run the async streaming generator
        asyncio.run(chat_stream())

        # Save assistant response
        st.session_state.history.append({"role": "assistant", "text": response_text})
```

```
...
Multi-agent AI system using Strands Agents to handle cloud cost analysis and infrastructure-as-code generation.
The system has three specialized agents a **Cost Explorer Agent**, a **Terraform Generator Agent**, and a **Supervisor (orchestrator) Agent**
A **Streamlit chatbot UI**.
...

from strands import Agent, tool
from strands.tools.mcp import MCPClient
from mcp.client.sse import sse_client

# (A) Cost Explorer Agent - uses MCP to fetch AWS cost data.
# -----
# Create an MCP client connected to the local Cost Explorer MCP server.
cost_mcp_client = MCPClient(lambda: sse_client("http://localhost:8080/sse"))
# List the tools provided by the MCP server (this fetches cost-related APIs)
cost_tools = []
with cost_mcp_client:
    cost_tools = cost_mcp_client.list_tools_sync() # list of Tool objects

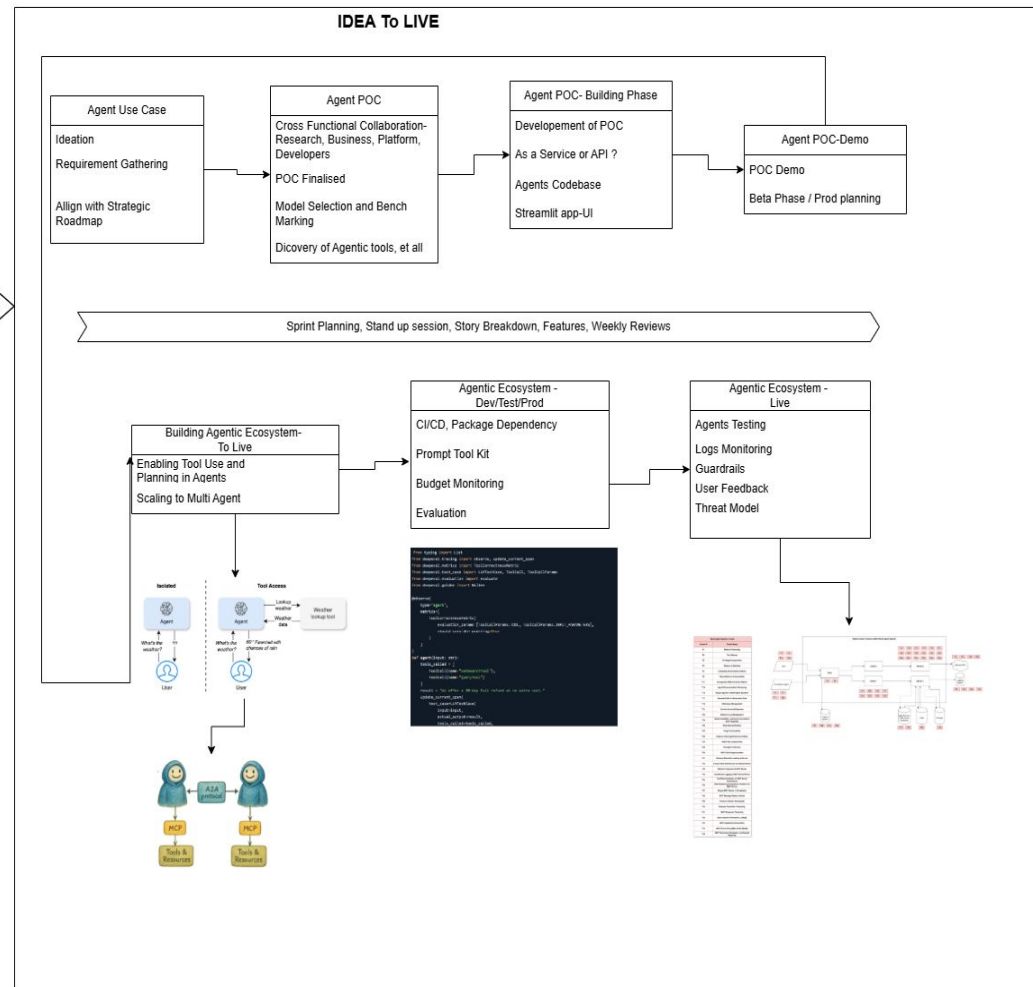
# Define a specialized cost-analysis agent as a tool.
COST_AGENT_PROMPT = (
    "You are a specialized AWS Cost Explorer assistant. "\n"
    "Given a user's request about AWS spending or cost breakdown, use the available cost tools "\n"
    "to query the AWS Cost Explorer data and return insights (service-wise spend, usage patterns, breakdowns, etc.) "\n"
    "in clear language."
)

@tool("cost_agent")
def cost_agent(query: str) -> str:
    try:
        # Create a new agent with the cost tools and a guiding prompt
        agent = Agent(
            system_prompt=COST_AGENT_PROMPT,
            tools=cost_tools,
            model = <Bedrock-Model>
        )
        response = agent(query)
        # The agent returns a textual answer including cost insights.
        return str(response)
    except Exception as e:
        return f"Error in Cost Agent: {str(e)}"

# (B) Terraform Generator Agent - generates Terraform code.
# -----
TERRAFORM_AGENT_PROMPT = (
    "You are an AWS Terraform expert. The user will describe an AWS architecture "\n"
    "or resource requirements in natural language. Generate a complete, valid Terraform "\n"
    "configuration (HCL code) that implements the described infrastructure. "\n"
    "Return only the Terraform code, with appropriate resources and modules."
)

@tool("terraform_agent")
def terraform_agent(query: str) -> str:
    try:
        agent = Agent(
            system_prompt=TERRAFORM_AGENT_PROMPT,
            tools=[] # no external tools needed, just rely on the LLM for code generation
            model = <Bedrock-Model>
        )
        response = agent(query)
        return str(response)
    except Exception as e:
        return f"Error in Terraform Agent: {str(e)}"
```

Idea to Live



Architecture Options




- Microservices for decoupling logic



- GenAI platforms (Bedrock, Azure AI Studio)



- Event-driven flows with queues (Pub/Sub, SQS)



- Stateless LLM 'brain' + API-based tools



- Recommended:
Modular, cloud-native,
agent-as-a-service

Key KPIs



- Task Efficiency – Completion rate, latency



- Reliability – Errors, uptime



- Engagement – Usage frequency, sessions



- Satisfaction – CSAT, feedback



- Business Impact – ROI, hours saved, cost/request

References

- [Agent Design Pattern Catalogue: A Collection of Architectural Patterns for Foundation Model based Agents](#)
- [Agentic or Tool use – Ragas](#)
- [Multi-agent Systems](#)
- [Introduction - Model Context Protocol](#)
- [ADK Tutorials! - Agent Development Kit](#)